Movie Script Summarization Experiments

Team members: Andy Fiegleman / Harry Lu / Stephen Tan

Reviewers: Mark, Peter

The overall score is the weighted sum of the individual components from the rubric. Note that most projects receive "*Satisfactory (83%-88%)*" for most components. The median overall grade for all projects this semester is 90.1%.

Overall Score: 91.35

- Overall Themes (10%): 94.0
- Crisp Objective (20%): 91.0
- Methodology & Analysis (50%): 90.5
- Technical Communication (20%): 92.5

The authors explore the challenging task of summarizing movie scripts which forces them to deal with long inputs. They adopt a variety of techniques to reduce the size of the input. They end up producing some summaries but provide a number of future paths to explore especially if not constrained by resources.

The overall theme is relatively unique and they provide a clear statement of the problem and their goals.

The objective is crisp and the metric is correctly applied. They mention looking at all ROUGE scores but ultimately only report ROUGE-1 which is just looking at unigrams. They manually evaluate a number of their output summaries which is a decent way to test readability and faithfulness but clearly isn't scalable.

The experimental design is quite sound. They identify a number of facets to the problem and proceed to systematically test them.

The authors provide some decent analysis of why their models perform as they do.

The paper is well structured and the presentation was good. All in all a very good job.